



dual1.0: phân bố latency mỗi predictor call — CPU OV async vs GPU sync
CPU OV chạy SONG SONG với GPU work → KHÔNG pause Llama dù chậm hơn


r = 2 req/s




r = 4 req/s




r = 8 req/s



r = 16 req/s



r = 32 req/s



r = 64 req/s

