




Phân bố ITL: opt-xxx có 'block tail' do predictor sync, dual1.0 không  
Đường nét đứt = baseline decode (median dual). Tỷ lệ ITL > 2×baseline cho thấy GPU bị block bao nhiêu thường xuyên.

**r = 2 req/s · P(ITL > 2×baseline): opt-xxx 7.9% vs dual1.0 6.7%** · **r = 4 req/s · P(ITL > 2×baseline): opt-xxx 12.3% vs dual1.0 8.6%** · **r = 8 req/s · P(ITL > 2×baseline): opt-xxx 16.5% vs dual 13.8%**



**r = 16 req/s · P(ITL > 2×baseline): opt-xxx 16.6% vs dual1.0 15.5%** · **r = 32 req/s · P(ITL > 2×baseline): opt-xxx 18.2% vs dual1.0 16.7%** · **r = 64 req/s · P(ITL > 2×baseline): opt-xxx 16.8% vs dual 15.4%**

